

SuperAI — Day 2 Debrief

Singapore · June 11, 2026

If Day 1 was the macro picture — capital cycles and the commoditization of the model layer — Day 2 was about where intelligence actually lands: finance, bio, health wearables, Southeast Asian physical AI, inference infrastructure, and the agent data infrastructure underneath. The panel lineup itself was an answer to the question everyone keeps asking: *does the app layer actually work?*

Three threads ran through the entire day.

Three Threads

1. Trust and governance are the rate-limiting step of deployment. Human-in-the-loop standards in regulated finance, "Know Your Agent" frameworks, CIO anxiety over shadow AI, and the diagnosis that intra-org trust is a prerequisite for agent collaboration showed up in nearly every session. The consensus: the bottleneck is trust, not technology. The mental model that kept recurring — onboard an agent like a new hire. Zero access on day one, access earned incrementally.

2. The web is becoming more hostile to agents. Smarter CAPTCHAs, honeypots, "pits and mazes" that trap crawlers in endless gibberish to drain compute, and fake data served to detected bots. Anti-automation investment is projected to more than triple by the end of the decade. The blunt takeaway: "looking human is the only viable strategy at scale" — and the open agentic web isn't here yet, so build for the web as it actually exists.

3. The moat is shifting from intelligence to unique data and physical assets. "Proprietary experimental data matters more than model architecture" (bio). "The real bottleneck isn't model intelligence, it's physical-world ground truth" (Grab). "Assembling the smartest people is no longer a moat" (EXO). As intelligence commoditizes, defensibility moves to what can't be copied: data, physical footprint, and product integration.

The single most direct market signal of the day: the **AI-native + Asian-footprint** configuration was named the highest-probability-of-success setup in the bio session.

Session Highlights

Inference & Scaling (Cerebras / ARM / Friendly AI). Inference — not training — is now the central challenge. The three players diagnosed the bottleneck differently: speed (Cerebras), the whole system rather than any single component (ARM), and raw GPU capacity plus reliability at scale (Friendly AI). The ecosystem is bifurcating into model labs (training) and inference providers (serving at scale). Cost per token has fallen 90%+ in three years while token demand is projected to grow ~24x by 2030, and pricing is shifting from subscription to usage-based. Open-source models are closing the quality gap at lower cost. Two-year outlook: new hardware enables fundamentally new applications, not just faster versions of old ones.

Southeast Asia & Physical AI (Grab). A three-wave framing: chatbots → agentic AI (now) → physical-world AI (next). The argument: the real constraint is the lack of physical-world ground truth. Concrete perception-layer products turn drivers into a distributed sensing layer (computer vision for rain/flood detection and rerouting), consolidate merchant CCTV into a "virtual store manager," and run on-device voice and safety AI across eight countries and dozens of languages — augmenting, not replacing, human supply.

Health Wearables (Oura / Superpower / Luki / Heidi). The thesis: data → insight → action, where wearables only matter when the AI acts proactively. Symptom detection 2–3 days before clinical onset; dedicated women's-health models; prevention-first care. The notable framing is the "health OS" — the long-run survivors will be horizontal aggregators connecting lab data, biometrics, and life context, agnostic to device brand. Five-year view: rising Asian wearable penetration and clinical-outcomes tracking as the unlock for health-system adoption.

Models & Cloud Strategy (Alibaba Cloud / Meitu / Video Rebirth). Why open source matters: it enables fine-tuning and inference on proprietary domain data, broadly accelerating enterprise adoption. Video/image models differ from LLMs — pixels and frames rather than tokens, with cross-frame consistency as the hard part — and have converged on a transformer-diffusion hybrid architecture. World models that understand and recreate the physical world are emerging for industrial and health use cases.

Productivity & Trust in Teams (EXO and others). Organizations are flattening; smaller teams doing parallel work is a key productivity lever; intra-org trust is a prerequisite for agents to work. Shadow AI is the top CIO concern, and there's no purely technical fix — it requires board-level governance and a clear policy on which agents touch which systems. A progressive trust model (start at zero access, reveal incrementally) and cloud-hosted agents (no company data by default, full audit visibility) were recurring design choices. On defensibility: embedding AI into the product — not just the process — is the fastest path to compounding advantage.

Agent Web Access (Oxylabs). The agentic web is still aspirational — protocols exist, meaningful adoption doesn't, and the underlying web is getting shakier. Drivers: scraping increasingly seen as adversarial, zero-click search cutting publisher revenue, e-commerce platforms restricting agent access, and fast-growing anti-automation spend. The only viable strategy at scale is to look human across behavior, passive fingerprinting, and browser-engine signals.

Data Infrastructure for Agents (Bright Data / Cohere / Clickhouse). "Real-time" has two dimensions: latency (UX) and freshness (output quality). The recommended pattern is a hybrid — a solid index for the majority of queries, augmented with real-time sourcing for time-sensitive ones. Quality = freshness + relevance; fresh data that returns the wrong result is still a failure. In regulated industries the primary challenge is trust and governance (RBAC, tool permissions, inherited user permissions), and accountability for bad agent decisions sits with the humans who design them — treat agent deployment like dev → staging → production.

Personal notes from Day 2. Shared for anyone who couldn't make it in the room.